

Ass1: Query Implementation using RDBM System and Elastic Search

(Issue: Jan 09, Due: Jan 25)

- TA: Trishla Shah trishla@dal.ca
- Ass1-Tut1: Jan 15, 4:05-5:25PM; Ass1-Tut2: Jan 22, 4:05-5:25PM, Chemistry 226
- Ass1-help Hours: Jan 11, 18, 24, 2:00-3:30PM, CS 233

1. Objectives:

- 1) To review conventional Relational Database Systems
- 2) To learn Infrastructure Services on a Cloud System
- 3) To learn Distributed Database concepts and implementation
- 4) To learn installation of Relational Database on a Cloud VM

2. Tasks:

- 1) Open a personal account on IBM Bluemix cloud service. (* If you prefer to use any other option, a justification must be provided.)
- 2) Create a DB instance running IBM's DBMS (DB2) on Bluemix cloud service.
 - a. You are given a dataset containing data instances of "Bus Routes", "Bus Stops" and "Bus Stop Times".
 - b. Create three tables : Trips, Stops and Stoptimes and Load data into the tables from the given data .csv files.
- 3) Perform following queries:
 - a. Find all buses for a particular Bus Stop
 1. Input: Bus Stop Name
 2. Output: List of all buses, response time for the search query
 - b. Find buses between two time ranges
 1. Input: Time Range 1 (hh:mm:ss), Time Range 2 (hh:mm:ss)
 2. Output: List of all buses, response time for the search query
 - c. List all the bus-stops in sequence of a particular bus on a particular route (Herein, route information signifies all the stops and their sequence made by the bus running on the provided route). (Hint: For particular route use route_id and for particular bus use trip_headsign)
 1. Input: Bus Name, Route Name
 2. Output: List of all routes, response time for the search query
 - d. Find top 3 bus stops that are the busiest throughout the day. (Hint: The bus stops with high volume of bus routes).
 1. Input: None
 2. Output: List of Bus Name, response time for the search query

Note: - During loading set maximum no. of warnings to 999999999 instead of 1000. This will allow to load the database upto the maximum warnings.
- Some rows in the database might get rejected due to noise or other errors present. That must be ignored.

- 4) Use IBM Bluemix's Compose for Elasticsearch database. Import data for Bus Route.
Run search query on all three documents.

Note: It is not necessary to use nested or join queries for elasticsearch. You can write separate sub queries and use hardcoded input (Hint: Output of first sub query will be become input for second sub query and so on.)

- 5) Observe the difference in performance and search quality. Explain why one would use Elastic search over the local search or vice versa.
- 6) Write a report including the following sections:
 - a. Task Description: Present the business scenario (i.e., the application and the requirements), and the DB (provide an instance of the database, and the reference source, etc.).
 - b. Relational Database Design: Provide an overview of schema for the database. Please mention Primary Key, Foreign Keys and Indexes if any. You can present this in simple text or an ERD Diagram.
 - c. Application Queries & Test Results: Please provide details on experimentation performed on relational database and elastic search. Please mention your search queries and response time.
 - d. Summary: Provide a summary of your work & observations on the application (i.e. the original DB) and the experience of using the software tools (i.e. your comments & recommendations, etc.).

3. Submit your Ass1 report electronically:

- 1) Please use Bright Space to submit your assignment
- 2) In addition to the report please also submit queries.
- 3) Please also provide any scripts that you used for implementation of Relational Database on Cloud and/or local system (if any).
- 4) Submit all items within one Zip file

*** Plagiarism and Intellectual Honesty:** (<http://plagiarism.dal.ca>)

Dalhousie University defines "plagiarism as the presentation of the work of another author in such a way as to give one's reader reason to think it to be one's own."

Plagiarism is considered a serious academic offense which may lead to loss of credit, suspension or expulsion from the University, or even the revocation of a degree.